Google Cloud

Course summary

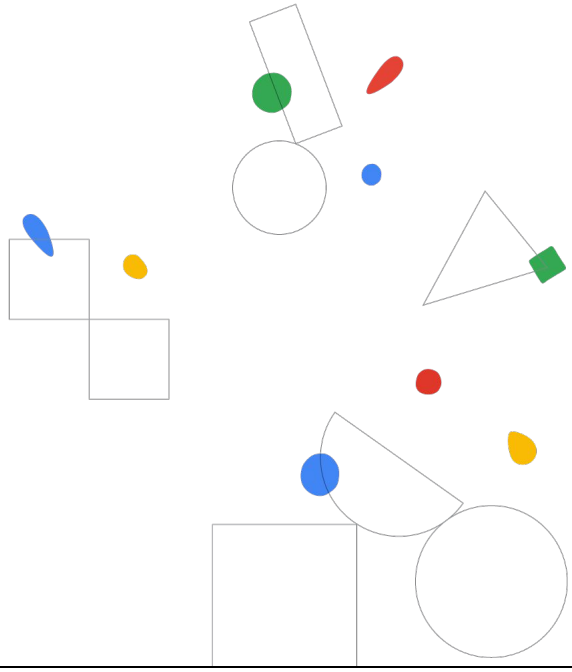

Production ML systems

This brings us to the end of the production ML systems course.

The first course in the specialization.

Before you go, let's quickly recap what you've learned.

01	Architecting Production ML Systems
02	Designing Adaptable ML Systems
03	Designing High-Performance ML Systems
04	Building Hybrid ML Systems

- In the first module, architecting production ML systems. We explored what an ML system should be able to do and the components that take responsibility for those actions. We also introduced two decisions that system architects will have to make whether to conduct dynamic or static training or even conduct dynamic or static inference.
- In module two, designing adaptable ML systems, you saw how change can affect an ML system and what can be done to mitigate those effects.
- In module three, designing high performance ML systems, we explored how to optimize the performance of an ML system by choosing the right hardware and removing bottlenecks.
- And finally, in module four, building hybrid ML systems, you learned about the technology behind hybrid systems that allows you to run your workloads on the cloud, on the edge using mobile devices or on premises.

Specialization

Advanced Machine Learning on Google Cloud

- Production Machine Learning Systems
- Image Processing and Generation with Google Cloud
- Sequence Models for Time Series and Natural Language Processing
- Recommendation Systems with TensorFlow on Google Cloud

We encourage you to continue to the next course …

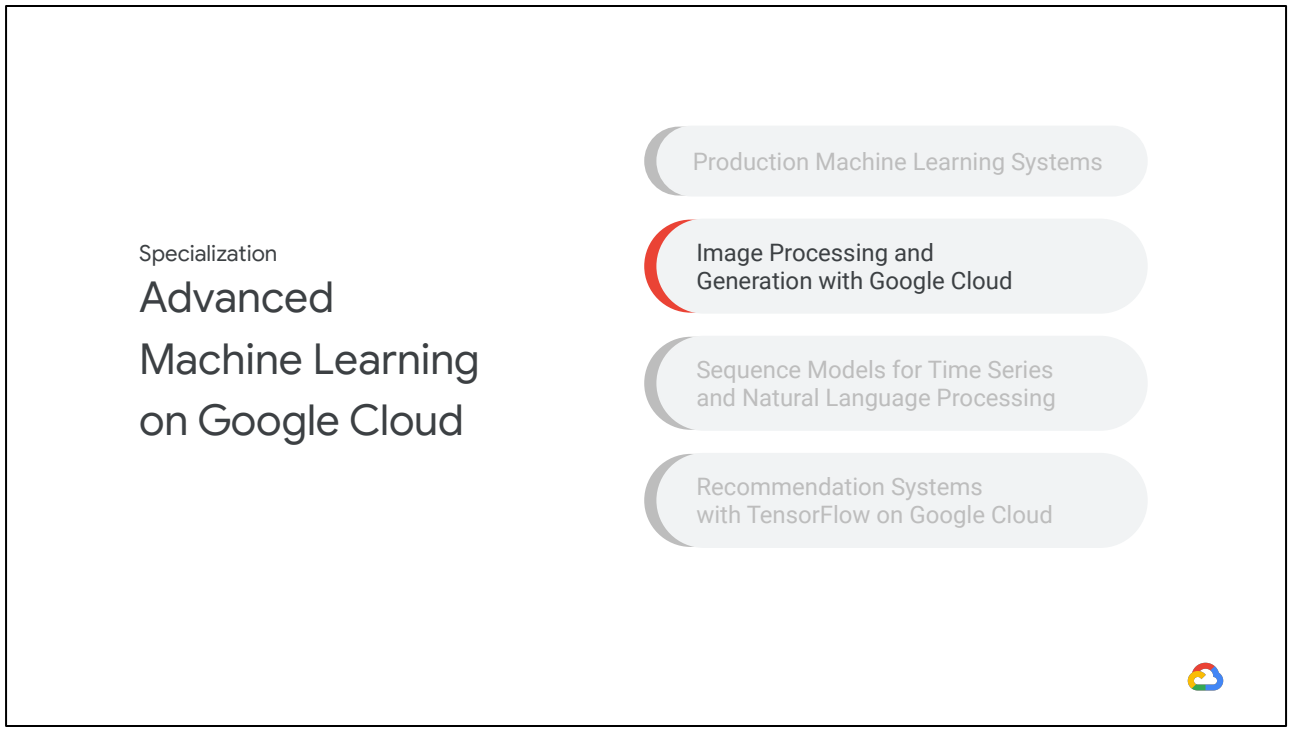

…image processing and generation with Google Cloud, where we'll explore convolutional networks, transfer learning and tensor processing units.

Thanks for learning with us.